

Text Mining with R *

Yanchang Zhao

<http://www.RDataMining.com>

R and Data Mining Course
Canberra, Australia

13 December 2018

*Chapter 10: Text Mining, in *R and Data Mining: Examples and Case Studies*.

Text Mining

Concept

Tasks

Twitter Data Analysis with R

Twitter

Extracting Tweets

Text Cleaning

Frequent Words and Word Cloud

Word Associations

Clustering

Topic Modelling

Sentiment Analysis

Follower Analysis

Retweeting Analysis

R Packages

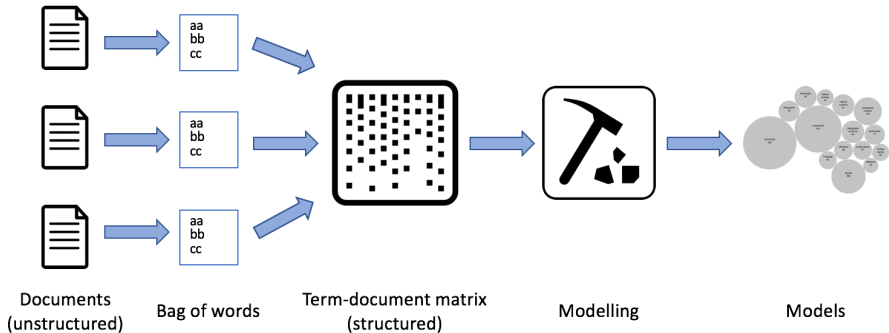
Wrap Up

Further Readings and Online Resources

- ▶ Text documents in a natural language
- ▶ Unstructured
- ▶ Documents in plain text, Word or PDF format
- ▶ Emails, online chat logs and phone transcripts
- ▶ Online news and forums, blogs, micro-blogs and social media
- ▶ ...

1. Transform text into structured data
 - ▶ Term-Document Matrix (TDM)
 - ▶ Entities and relations
 - ▶ ...
2. Apply traditional data mining techniques to the above structured data
 - ▶ Clustering
 - ▶ Classification
 - ▶ Social Network Analysis (SNA)
 - ▶ ...

Typical Process of Text Mining (cont.)



- ▶ Also known as Document-Term Matrix (DTM)
- ▶ A 2D matrix
- ▶ Rows: terms or words
- ▶ Columns: documents
- ▶ Entry $m_{i,j}$: number of occurrences of term t_i in document d_j
- ▶ Term weighting schemes: Term Frequency, Binary Weight, TF-IDF, etc.

- ▶ Term Frequency (TF) $tf_{i,j}$: the number of occurrences of term t_i in document d_j
- ▶ Inverse Document Frequency (IDF) for term t_i is:

$$idf_i = \log_2 \frac{|D|}{|\{d \mid t_i \in d\}|} \quad (1)$$

$|D|$: the total number of documents

$|\{d \mid t_i \in d\}|$: the number of documents where term t_i appears

- ▶ Term Frequency - Inverse Document Frequency (TF-IDF)

$$tfidf = tf_{i,j} \cdot idf_i \quad (2)$$

- ▶ IDF reduces the weight of terms that occur frequently in documents and increases the weight of terms that occur rarely.

An Example of TDM

Doc1: I like R.

Doc2: I like Python.

Term Frequency

	Doc1	Doc2
I	1	1
like	1	1
Python	0	1
R	1	0

IDF

	IDF
I	0
like	0
Python	1
R	1

TF-IDF

	Doc1	Doc2
I	0	0
like	0	0
Python	0	1
R	1	0

An Example of TDM

Doc1: I like R.

Doc2: I like Python.

Term Frequency

	Doc1	Doc2
I	1	1
like	1	1
Python	0	1
R	1	0

IDF

	IDF
I	0
like	0
Python	1
R	1

TF-IDF

	Doc1	Doc2
I	0	0
like	0	0
Python	0	1
R	1	0

Terms that can distinguish different documents are given greater weights.

An Example of TDM (cont.)

Doc1: I like R.

Doc2: I like Python.

Term Frequency

	Doc1	Doc2
I	1	1
like	1	1
Python	0	1
R	1	0

IDF

	IDF
I	0
like	0
Python	1
R	1

Normalized Term Frequency

	Doc1	Doc2
I	0.33	0.33
like	0.33	0.33
Python	0	0.33
R	0.33	0

Normalized TF-IDF

	Doc1	Doc2
I	0	0
like	0	0
Python	0	0.33
R	0.33	0

An Example of Term Weighting in R

```
library(magrittr)
library(tm) ## package for text mining
a <- c("I like R", "I like Python")
## build corpus
b <- a %>% VectorSource() %>% Corpus()
## build term document matrix
m <- b %>% TermDocumentMatrix(control=list(wordLengths=c(1, Inf)))
m %>% inspect()
## various term weighting schemes
m %>% weightBin() %>% inspect() ## binary weighting
m %>% weightTf() %>% inspect() ## term frequency
m %>% weightTfIdf(normalize=F) %>% inspect() ## TF-IDF
m %>% weightTfIdf(normalize=T) %>% inspect() ## normalized TF-IDF
```

More options provided in package *tm*:

- ▶ weightSMART
- ▶ WeightFunction

- ▶ Text classification
- ▶ Text clustering and categorization
- ▶ Topic modelling
- ▶ Sentiment analysis
- ▶ Document summarization
- ▶ Entity and relation extraction
- ▶ ...

- ▶ To identify topics in a set of documents
- ▶ It groups both documents that use similar words and words that occur in a similar set of documents.
- ▶ Intuition: Documents related to R would contain more words like R, ggplot2, plyr, stringr, knitr and other R packages, than Python related keywords like Python, NumPy, SciPy, Matplotlib, etc.
- ▶ A document can be of multiple topics in different proportions. For instance, a document can be 90% about R and 10% about Python. \Rightarrow soft/fuzzy clustering
- ▶ Latent Dirichlet Allocation (LDA): the most widely used topic model


- ▶ Also known as opinion mining
- ▶ To determine attitude, polarity or emotions from documents
- ▶ Polarity: positive, negative, netural
- ▶ Emotions: angry, sad, happy, bored, afraid, etc.
- ▶ Method:
 1. identify invidual words and phrases and map them to different emotional scales
 2. adjust the sentiment value of a concept based on modifications surrounding it

- ▶ To create a summary with major points of the original document
- ▶ Approaches
 - ▶ Extraction: select a subset of existing words, phrases or sentences to build a summary
 - ▶ Abstraction: use natural language generation techniques to build a summary that is similar to natural language

- ▶ Named Entity Recognition (NER): identify named entities in text into pre-defined categories, such as person names, organizations, locations, date and time, etc.
- ▶ Relationship Extraction: identify associations among entities
- ▶ Example:
Ben lives at 5 Geroqe St, Sydney.

- ▶ Named Entity Recognition (NER): identify named entities in text into pre-defined categories, such as person names, organizations, locations, date and time, etc.
- ▶ Relationship Extraction: identify associations among entities
- ▶ Example:
Ben lives at 5 Geroge St, Sydney.

- ▶ Named Entity Recognition (NER): identify named entities in text into pre-defined categories, such as person names, organizations, locations, date and time, etc.
- ▶ Relationship Extraction: identify associations among entities
- ▶ Example:
Ben lives at 5 Geroge St, Sydney.



Text Mining

Concept

Tasks

Twitter Data Analysis with R

Twitter

Extracting Tweets

Text Cleaning

Frequent Words and Word Cloud

Word Associations

Clustering

Topic Modelling

Sentiment Analysis

Follower Analysis

Retweeting Analysis

R Packages

Wrap Up

Further Readings and Online Resources



- ▶ An online social networking service that enables users to send and read short 280-character (used to be 140 before November 2017) messages called “tweets” (Wikipedia)
- ▶ Over 300 million monthly active users (as of 2018)
- ▶ Creating over 500 million tweets per day

RDataMining Twitter Account



Yanchang Zhao

@RDataMining

Tweets

748

Following

85

Followers

3,332

Likes

17

Lists

0

Moments

0

Edit profile

@RDataMining

R and Data Mining. Group on LinkedIn:
group.rdatamining.com

Australia

RDataMining.com

Joined April 2011

Photos and videos



Your Tweet activity

Your Tweets earned **3,421 impressions**
over the last **28 days**



Yanchang Zhao @RDataMining · Nov 26

I will deliver a tutorial on R and Data Mining at the AusDM 2018 Conference, in Bathurst, Australia on 29 November 2018. Slides, script and data for the tutorial are available at rdatamining.com/training/ausdm...



2



2



Yanchang Zhao @RDataMining · Nov 14

New webinar: 10 Guidelines for A/B Testing

Discover 10 best practices that will help you avoid common A/B testing pitfalls with Emily Robinson, a data scientist at DataCamp.

Registration link: attendee.gotowebinar.com/register/84237...



1



Yanchang Zhao @RDataMining · Nov 7

Research Scientist in Data Science - 2 Positions at Data61, CSIRO. Application close on 28 November 2018. jobs.csiro.au/job/Melbourne%...



1



1. Extract tweets and followers from the Twitter website with R and the *twitteR* package
2. With the *tm* package, clean text by removing punctuations, numbers, hyperlinks and stop words, followed by stemming and stem completion
3. Build a term-document matrix
4. Cluster Tweets with text clustering
5. Analyse topics with the *topicmodels* package
6. Analyse sentiment with the *sentiment140* package
7. Analyse following/followed and retweeting relationships with the *igraph* package

```
## Option 1: retrieve tweets from Twitter
library(twitteR)
library(ROAuth)
## Twitter authentication
setup_twitter_oauth(consumer_key, consumer_secret, access_token, access_secret)
## 3200 is the maximum to retrieve
tweets <- "RDataMining" %>% userTimeline(n = 3200)
```

See details of *Twitter Authentication with OAuth* in Section 3 of <http://geoffjentry.hexdump.org/twitterR.pdf>.

```
## Option 2: download @RDataMining tweets from RDataMining.com
library(twitteR)
url <- "http://www.rdatamining.com/data/RDataMining-Tweets-20160212.rds"
download.file(url, destfile = "./data/RDataMining-Tweets-20160212.rds")
## load tweets into R
tweets <- readRDS("./data/RDataMining-Tweets-20160212.rds")
```

```
(n.tweet <- tweets %>% length())
## [1] 448

# convert tweets to a data frame
tweets.df <- tweets %>% twListToDF()
# tweet #1
tweets.df[1, c("id", "created", "screenName", "replyToSN",
  "favoriteCount", "retweetCount", "longitude", "latitude", "text")]
##              id          created screenName replyToSN
## 1 697031245503418368 2016-02-09 12:16:13 RDataMining    <NA>
##   favoriteCount retweetCount longitude latitude
## 1             13             14         NA         NA
##                                     ...
## 1 A Twitter dataset for text mining: @RDataMining Tweets ex...

# print tweet #1 and make text fit for slide width
tweets.df$text[1] %>% strwrap(60) %>% writeLines()
## A Twitter dataset for text mining: @RDataMining Tweets
## extracted on 3 February 2016. Download it at
## https://t.co/lQp94IvfPf
```


- ▶ Convert to lower case: `tolower`
- ▶ Remove punctuation: `removePunctuation`
- ▶ Remove numbers: `removeNumbers`
- ▶ Remove URLs
- ▶ Remove stop words (like 'a', 'the', 'in'): `removeWords`, `stopwords`
- ▶ Remove extra white space: `stripWhitespace`

```
library(tm)
# function for removing URLs, i.e.,
# "http" followed by any non-space letters
removeURL <- function(x) gsub("http[^[[:space:]]*", "", x)
# function for removing anything other than English letters or space
removeNumPunct <- function(x) gsub("[^[:alpha:][:space:]]*", "", x)
# customize stop words
myStopwords <- c(setdiff(stopwords('english'), c("r", "big")),
                  "use", "see", "used", "via", "amp")
```

See details of regular expressions by running `?regex` in R console.

```
# build a corpus and specify the source to be character vectors
corpus.raw <- tweets.df$text %>% VectorSource() %>% Corpus()

# text cleaning
corpus.cleaned <- corpus.raw %>%
  # convert to lower case
  tm_map(content_transformer(tolower)) %>%
  # remove URLs
  tm_map(content_transformer(removeURL)) %>%
  # remove numbers and punctuations
  tm_map(content_transformer(removeNumPunct)) %>%
  # remove stopwords
  tm_map(removeWords, myStopwords) %>%
  # remove extra whitespace
  tm_map(stripWhitespace)
```

```
## stem words
corpus.stemmed <- corpus.cleaned %>% tm_map(stemDocument)

## stem completion
stemCompletion2 <- function(x, dictionary) {
  x <- unlist(strsplit(as.character(x), " "))
  x <- x[x != ""]
  x <- stemCompletion(x, dictionary=dictionary)
  x <- paste(x, sep=" ", collapse=" ")
  stripWhitespace(x)
}

corpus.completed <- corpus.stemmed %>%
  lapply(stemCompletion2, dictionary=corpus.cleaned) %>%
  VectorSource() %>% Corpus()
```

[†] <http://stackoverflow.com/questions/25206049/stemcompletion-is-not-working>

Before/After Text Cleaning and Stemming

original text

```
corpus.raw[[1]]$content %>% strwrap(60) %>% writeLines()  
## A Twitter dataset for text mining: @RDataMining Tweets  
## extracted on 3 February 2016. Download it at  
## https://t.co/lQp94IvfPf
```

after basic cleaning

```
corpus.cleaned[[1]]$content %>% strwrap(60) %>% writeLines()  
## twitter dataset text mining rdatamining tweets extracted  
## february download
```

stemmed text

```
corpus.stemmed[[1]]$content %>% strwrap(60) %>% writeLines()  
## twitter dataset text mine rdatamin tweet extract february  
## download
```

after stem completion

```
corpus.completed[[1]]$content %>% strwrap(60) %>% writeLines()  
## twitter dataset text miner rdatamining tweet extract  
## download
```

Issues in Stem Completion: “Miner” vs “Mining”

```
# count word frequency
wordFreq <- function(corpus, word) {
  results <- lapply(corpus,
    function(x) grep(as.character(x), pattern=paste0("\\<",word)) )
  sum(unlist(results))
}
n.miner <- corpus.cleaned %>% wordFreq("miner")
n.mining <- corpus.cleaned %>% wordFreq("mining")
cat(n.miner, n.mining)
## 9 104


# replace old word with new word
replaceWord <- function(corpus, oldword, newword) {
  tm_map(corpus, content_transformer(gsub),
    pattern=oldword, replacement=newword)
}
corpus.completed <- corpus.completed %>%
  replaceWord("miner", "mining") %>%
  replaceWord("universidad", "university") %>%
  replaceWord("scienc", "science")
```

```
tdm <- corpus.completed %>%  
  TermDocumentMatrix(control = list(wordLengths = c(1, Inf))) %>%  
  print  
  
## <<TermDocumentMatrix (terms: 1073, documents: 448)>>  
## Non-/sparse entries: 3594/477110  
## Sparsity          : 99%  
## Maximal term length: 23  
## Weighting          : term frequency (tf)  
  
idx <- which(dimnames(tdm)$Terms %in% c("r", "data", "mining"))  
tdm[idx, 21:30] %>% as.matrix()  
  
##           Docs  
## Terms      21 22 23 24 25 26 27 28 29 30  
## mining     0  0  0  0  1  0  0  0  0  1  
## data       0  1  0  0  1  0  0  0  0  1  
## r          1  1  1  1  0  1  0  1  1  1
```

```
# inspect frequent words
freq.terms <- tdm %>% findFreqTerms(lowfreq = 20) %>% print
## [1] "mining"      "rdatamining" "text"      "analytics"
## [5] "australia"   "data"        "canberra"  "group"
## [9] "university"  "science"     "slide"     "tutorial"
## [13] "big"         "learn"       "package"   "r"
## [17] "network"     "course"      "introduction" "talk"
## [21] "analysing"   "research"    "position"  "example"

term.freq <- tdm %>% as.matrix() %>% rowSums()
term.freq <- term.freq %>% subset(term.freq >= 20)
df <- data.frame(term = names(term.freq), freq = term.freq)
```

```
library(ggplot2)
ggplot(df, aes(x=term, y=freq)) + geom_bar(stat="identity") +
  xlab("Terms") + ylab("Count") + coord_flip() +
  theme(axis.text=element_text(size=7))
```




```
m <- tdm %>% as.matrix
# calculate the frequency of words and sort it by frequency
word.freq <- m %>% rowSums() %>% sort(decreasing = T)
# colors
library(RColorBrewer)
pal <- brewer.pal(9, "BuGn")[-(1:4)]
```

```
# plot word cloud
library(wordcloud)
wordcloud(words = names(word.freq), freq = word.freq, min.freq = 3,
  random.order = F, colors = pal)
```


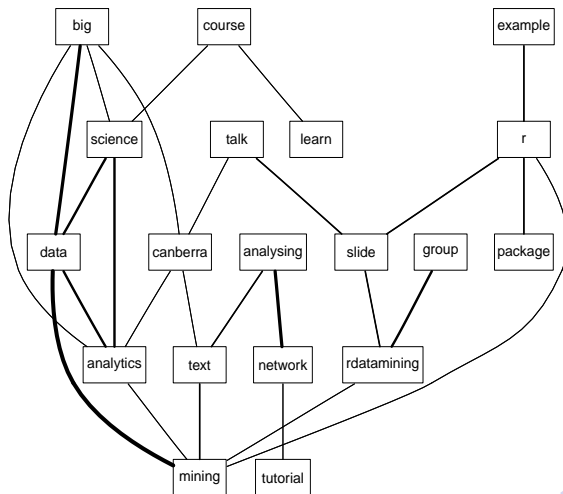


```
# which words are associated with 'r'?
tdm %>% findAssocs("r", 0.2)
## $r
##      code  example  series  user markdown
##      0.27    0.21    0.21    0.20    0.20
```

```
# which words are associated with 'data'?
tdm %>% findAssocs("data", 0.2)
## $data
##      mining      big analytics  science  poll
##      0.48      0.44      0.31      0.29  0.24
```

Network of Terms

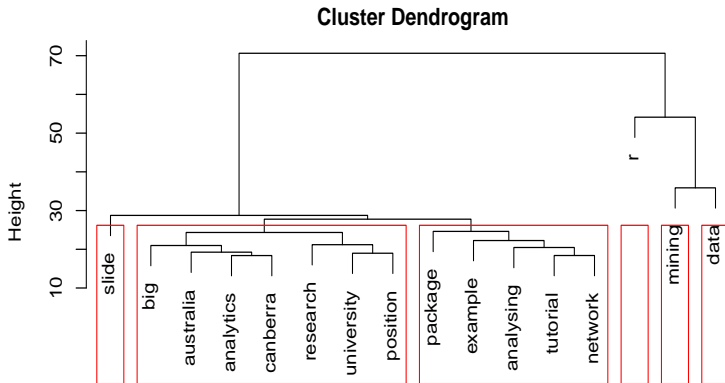
```
library(graph)
library(Rgraphviz)
plot(tdm, term = freq.terms, corThreshold = 0.1, weighting = T)
```



Hierarchical Clustering of Terms

```
# remove sparse terms
m2 <- tdm %>% removeSparseTerms(sparse = 0.95) %>% as.matrix()
# calculate distance matrix
dist.matrix <- m2 %>% scale() %>% dist()
# hierarchical clustering
fit <- dist.matrix %>% hclust(method = "ward")
```

```
plot(fit)
fit %>% rect.hclust(k = 6)  # cut tree into 6 clusters
groups <- fit %>% cutree(k = 6)
```



hclust (*, "ward.D")

```

m3 <- m2 %>% t() # transpose the matrix to cluster documents
set.seed(122) # set a fixed random seed to make the result reproducibl
k <- 6 # number of clusters
kmeansResult <- kmeans(m3, k)
round(kmeansResult$centers, digits = 3) # cluster centers

```

##	mining	analytics	australia	data	canberra	university	slide
## 1	0.435	0.000	0.000	0.217	0.000	0.000	0.087
## 2	1.128	0.154	0.000	1.333	0.026	0.051	0.179
## 3	0.055	0.018	0.009	0.164	0.027	0.009	0.227
## 4	0.083	0.014	0.056	0.000	0.035	0.097	0.090
## 5	0.412	0.206	0.098	1.196	0.137	0.039	0.078
## 6	0.167	0.133	0.133	0.567	0.033	0.233	0.000

##	tutorial	big	package	r	network	analysing	research
## 1	0.043	0.000	0.043	1.130	0.087	0.174	0.000
## 2	0.026	0.077	0.282	1.103	0.000	0.051	0.000
## 3	0.064	0.018	0.109	1.127	0.045	0.109	0.000
## 4	0.056	0.007	0.090	0.000	0.090	0.111	0.000
## 5	0.059	0.333	0.010	0.020	0.020	0.059	0.020
## 6	0.000	0.167	0.033	0.000	0.067	0.100	1.233


##	position	example
## 1	0.000	1.043
## 2	0.000	0.026
## 3	0.000	0.000
## 4	0.076	0.035

```
for (i in 1:k) {  
  cat(paste("cluster ", i, ": ", sep = ""))  
  s <- sort(kmeansResult$centers[i, ], decreasing = T)  
  cat(names(s)[1:5], "\n")  
  # print the tweets of every cluster  
  # print(tweets[which(kmeansResult$cluster==i)])  
}  
  
## cluster 1:  r example mining data analysing  
## cluster 2:  data mining r package slide  
## cluster 3:  r slide data package analysing  
## cluster 4:  analysing university slide package network  
## cluster 5:  data mining big analytics canberra  
## cluster 6:  research data position university mining
```



```
dtm <- tdm %>% as.DocumentTermMatrix()
library(topicmodels)
lda <- LDA(dtm, k = 8) # find 8 topics
term <- terms(lda, 7) # first 7 terms of every topic
term <- apply(term, MARGIN = 2, paste, collapse = ", ") %>% print
## To...
## "r, mining, package, research, analytics, group, ...
## To...
## "data, r, mining, analytics, big, australia, s...
## To...
## "r, data, mining, slide, example, useful,...
## To...
## "data, r, mining, research, introduction, example, s...
## To...
## "data, mining, analysing, package, science, position, analy...
## To...
## "data, r, mining, slide, analysing, big, introduc...
## To...
## "big, data, ausdm, research, text, university, submis...
## To...
## "r, mining, data, slide, analysing, big, austr...
```

```
rdm.topics <- topics(lda) # 1st topic identified for every document (tw  
rdm.topics <- data.frame(date=as.IDate(tweets.df$created),  
                           topic=rdm.topics)  
ggplot(rdm.topics, aes(date, fill = term[topic])) +  
  geom_density(position = "stack")
```



Another way to plot steam graph:

```
# install package sentiment140
require(devtools)
install_github("sentiment140", "okugami79")
```

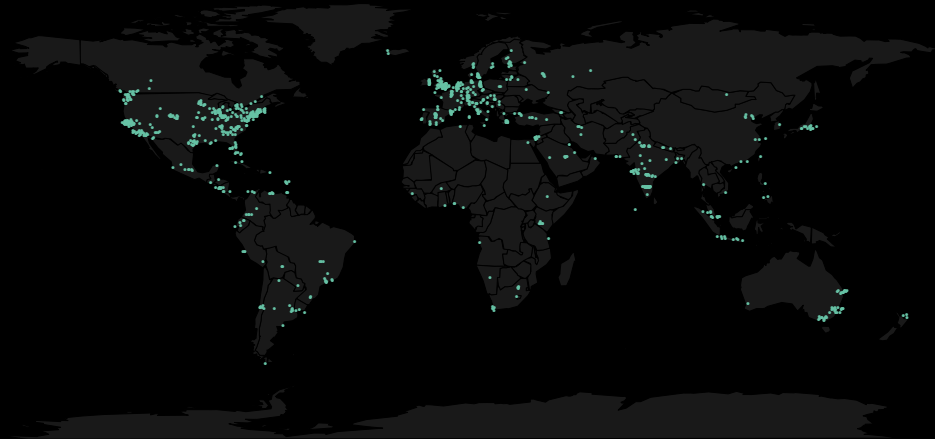
```
# sentiment analysis
library(sentiment)
sentiments <- sentiment(tweets.df$text)
table(sentiments$polarity)
# sentiment plot
sentiments$score <- 0
sentiments$score[sentiments$polarity == "positive"] <- 1
sentiments$score[sentiments$polarity == "negative"] <- -1
sentiments$date <- as.IDate(tweets.df$created)
result <- aggregate(score ~ date, data = sentiments, sum)
```

Retrieve User Info and Followers

```
user <- getUser("RDataMining")
user$toDataFrame()
friends <- user$getFriends() # who this user follows
followers <- user$getFollowers() # this user's followers
followers2 <- followers[[1]]$getFollowers() # a follower's followers
```


```
##           [,1] ...
## description "R and Data Mining. Group on LinkedIn: ht...
## statusesCount "583" ...
## followersCount "2376" ...
## favoritesCount "6" ...
## friendsCount "72" ...
## url "http://t.co/LwL50uRmPd" ...
## name "Yanchang Zhao" ...
## created "2011-04-04 09:15:43" ...
## protected "FALSE" ...
## verified "FALSE" ...
## screenName "RDataMining" ...
## location "Australia" ...
## lang "en" ...
## id "276895537" ...
## linkID "157" ...
```

@RDataMining Followers (#: 2376)



[‡]Based on Jeff Leek's twitterMap function at
<http://biostat.jhsph.edu/~jleek/code/twitterMap.R>

Active Influential Followers




Top Retweeted Tweets

```
# select top retweeted tweets
table(tweets.df$retweetCount)
selected <- which(tweets.df$retweetCount >= 9)

# plot them
dates <- strptime(tweets.df$created, format="%Y-%m-%d")
plot(x=dates, y=tweets.df$retweetCount, type="l", col="grey",
     xlab="Date", ylab="Times retweeted")
colors <- rainbow(10)[1:length(selected)]
points(dates[selected], tweets.df$retweetCount[selected],
       pch=19, col=colors)
text(dates[selected], tweets.df$retweetCount[selected],
     tweets.df$text[selected], col=colors, cex=.9)
```

Top Retweeted Tweets



Tracking Message Propagation

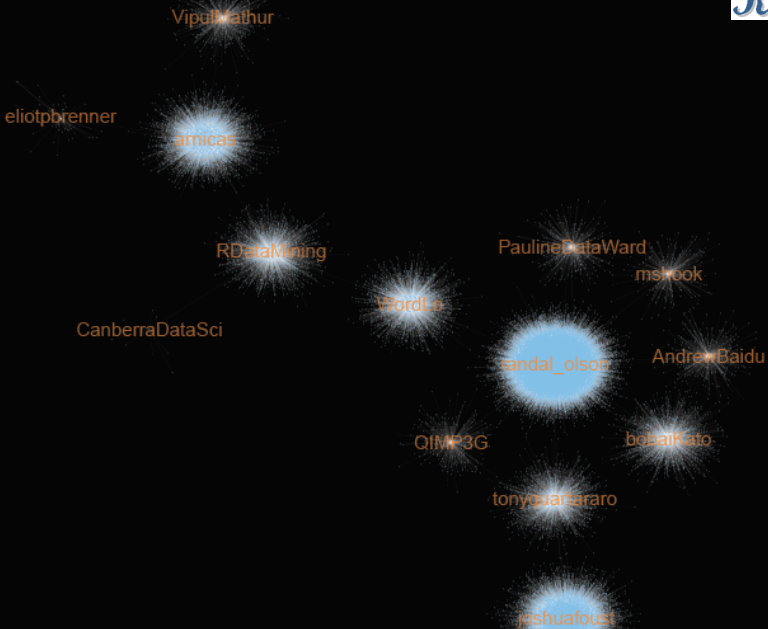
```
tweets[[1]]  
retweeters(tweets[[1]]$id)  
retweets(tweets[[1]]$id)
```

```
## [1] "RDataMining: A Twitter dataset for text mining: @RData..."
```

```
## [1] "197489286" "316875164" "229796464" "3316009302"  
## [5] "244077734" "16900353" "2404767650" "222061895"  
## [9] "11686382" "190569306" "49413866" "187048879"  
## [13] "6146692" "2591996912"
```

```
## [[1]]  
## [1] "bobaiKato: RT @RDataMining: A Twitter dataset for text..."  
##  
## [[2]]  
## [1] "VipulMathur: RT @RDataMining: A Twitter dataset for te..."  
##  
## [[3]]  
## [1] "tau_phoenix: RT @RDataMining: A Twitter dataset for te..."
```

The tweet potentially reached around 120,000 users.



Text Mining

Concept

Tasks

Twitter Data Analysis with R

Twitter

Extracting Tweets

Text Cleaning

Frequent Words and Word Cloud

Word Associations

Clustering

Topic Modelling

Sentiment Analysis

Follower Analysis

Retweeting Analysis

R Packages

Wrap Up

Further Readings and Online Resources

- ▶ Twitter data extraction: *twitteR*
- ▶ Text cleaning and mining: *tm*
- ▶ Word cloud: *wordcloud*
- ▶ Topic modelling: *topicmodels*, *lda*
- ▶ Sentiment analysis: *sentiment140*
- ▶ Social network analysis: *igraph*, *sna*
- ▶ Visualisation: *wordcloud*, *Rgraphviz*, *ggplot2*

- ▶ `userTimeline`, `homeTimeline`, `mentions`, `retweetsOfMe`: retrieve various timelines
- ▶ `getUser`, `lookupUsers`: get information of Twitter user(s)
- ▶ `getFollowers`, `getFollowerIDs`: retrieve followers (or their IDs)
- ▶ `getFriends`, `getFriendIDs`: return a list of Twitter users (or user IDs) that a user follows
- ▶ `retweets`, `retweeters`: return retweets or users who retweeted a tweet
- ▶ `searchTwitter`: issue a search of Twitter
- ▶ `getCurRateLimitInfo`: retrieve current rate limit information
- ▶ `twListToDF`: convert into `data.frame`

§ <https://cran.r-project.org/package=twitteR>

- ▶ `removeNumbers`, `removePunctuation`, `removeWords`, `removeSparseTerms`, `stripWhitespace`: remove numbers, punctuations, words or extra whitespaces
- ▶ `removeSparseTerms`: remove sparse terms from a term-document matrix
- ▶ `stopwords`: various kinds of stopwords
- ▶ `stemDocument`, `stemCompletion`: stem words and complete stems
- ▶ `TermDocumentMatrix`, `DocumentTermMatrix`: build a term-document matrix or a document-term matrix
- ▶ `termFreq`: generate a term frequency vector
- ▶ `findFreqTerms`, `findAssocs`: find frequent terms or associations of terms
- ▶ `weightBin`, `weightTf`, `weightTfIdf`, `weightSMART`, `WeightFunction`: various ways to weight a term-document matrix

Topic Modelling and Sentiment Analysis – Packages

topicmodels & *sentiment140*

Package *topicmodels* ||

- ▶ LDA: build a Latent Dirichlet Allocation (LDA) model
- ▶ CTM: build a Correlated Topic Model (CTM) model
- ▶ terms: extract the most likely terms for each topic
- ▶ topics: extract the most likely topics for each document

Package *sentiment140* **

- ▶ sentiment: sentiment analysis with the sentiment140 API, tune to Twitter text analysis

|| <https://cran.r-project.org/package=topicmodels>

** <https://github.com/okugami79/sentiment140>

Social Network Analysis and Visualization – Package

igraph ^{††}

- ▶ degree, betweenness, closeness, transitivity: various centrality scores
- ▶ neighborhood: neighborhood of graph vertices
- ▶ cliques, largest.cliques, maximal.cliques, clique.number: find cliques, ie. complete subgraphs
- ▶ clusters, no.clusters: maximal connected components of a graph and the number of them
- ▶ fastgreedy.community, spinglass.community: community detection
- ▶ cohesive.blocks: calculate cohesive blocks
- ▶ induced.subgraph: create a subgraph of a graph (*igraph*)
- ▶ read.graph, write.graph: read and writ graphs from and to files of various formats

^{††}<https://cran.r-project.org/package=igraph>

Text Mining

Concept

Tasks

Twitter Data Analysis with R

Twitter

Extracting Tweets

Text Cleaning

Frequent Words and Word Cloud

Word Associations

Clustering

Topic Modelling

Sentiment Analysis

Follower Analysis

Retweeting Analysis

R Packages

Wrap Up

Further Readings and Online Resources

- ▶ Transform unstructured data into structured data (i.e., term-document matrix), and then apply traditional data mining algorithms like clustering and classification
- ▶ Feature extraction: term frequency, TF-IDF and many others
- ▶ Text cleaning: lower case, removing numbers, punctuations and URLs, stop words, stemming and stem completion
- ▶ Stem completion may not always work as expected.
- ▶ Documents in languages other than English

Text Mining

Concept

Tasks

Twitter Data Analysis with R

Twitter

Extracting Tweets

Text Cleaning

Frequent Words and Word Cloud

Word Associations

Clustering

Topic Modelling

Sentiment Analysis

Follower Analysis

Retweeting Analysis

R Packages

Wrap Up

Further Readings and Online Resources

- ▶ Text Mining

https://en.wikipedia.org/wiki/Text_mining

- ▶ TF-IDF

<https://en.wikipedia.org/wiki/Tf\0T1\textendashidf>

- ▶ Topic Modelling

https://en.wikipedia.org/wiki/Topic_model

- ▶ Sentiment Analysis

https://en.wikipedia.org/wiki/Sentiment_analysis

- ▶ Document Summarization

https://en.wikipedia.org/wiki/Automatic_summarization

- ▶ Natural Language Processing

https://en.wikipedia.org/wiki/Natural_language_processing

- ▶ An introduction to text mining by Ian Witten

<http://www.cs.waikato.ac.nz/%7Eihw/papers/04-IHW-Textmining.pdf>

- ▶ Chapter 10 – Text Mining, in book *R and Data Mining: Examples and Case Studies*

<http://www.rdatamining.com/docs/RDataMining-book.pdf>

- ▶ RDataMining Reference Card

<http://www.rdatamining.com/docs/RDataMining-reference-card.pdf>

- ▶ Free online courses and documents

<http://www.rdatamining.com/resources/>



- ▶ RDataMining Group on LinkedIn (26,000+ members)

<http://group.rdatamining.com>

- ▶ Twitter (3,300+ followers)

@RDataMining

- ▶ Yanchang Zhao. *R and Data Mining: Examples and Case Studies*. ISBN 978-0-12-396963-7, December 2012. Academic Press, Elsevier. 256 pages.
<http://www.rdatamining.com/docs/RDataMining-book.pdf>
- ▶ Yanchang Zhao and Yonghua Cen (Eds.). *Data Mining Applications with R*. ISBN 978-0124115118, December 2013. Academic Press, Elsevier.
- ▶ Yanchang Zhao. Analysing Twitter Data with Text Mining and Social Network Analysis. In *Proc. of the 11th Australasian Data Mining & Analytics Conference (AusDM 2013)*, Canberra, Australia, November 13—15, 2013.



Thanks!

Email: [yanchang\(at\)RDataMining.com](mailto:yanchang(at)RDataMining.com)

Twitter: @RDataMining